

# OLEKSIY DYAGILEV

*13+ years of software engineering experience*

blog <http://dyagilev.org/>

github <https://github.com/fe2s>

I help companies to design and build software.

My interests are Big Data, distributed systems, stream processing and reliable systems.

Currently my primary skills are: **Scala, Spark and Kafka.**

## SKILLS

### Scala

- Akka-http
- Slick

### Big Data and NoSQL

- Apache Spark
- Apache Kafka
- Cassandra
- Redis
- Hadoop2 (YARN, HDFS, Hive)
- Apache Storm
- MongoDB
- Memcached
- Zookeeper
- GigaSpace XAP

### Java

- Java EE
- Spring MVC, Spring Data, Spring Session, Spring Cloud
- Netflix OSS(Hystrix, Feign)
- Maven

### Misc

- Linux
- AWS
- Docker
- Machine learning
- Algorithms

## WORK EXPERIENCE

*Dec-2017 - now, Upwork (contract)*

Designing and implementing real-time analytics application for trading market.

**Technologies:** Spark, Kafka, Cassandra, Scala, Akka-http, AWS S3

*Jul-2018 - now, RedisLabs (contract)*

Designing, developing and maintaining <https://github.com/RedisLabs/spark-redis> library. Implemented support of DataFrames and Structured Streaming.

**Technologies:** Spark, Redis, Scala

*Oct-2018 - Feb 2019, Chief Software Engineer, EPAM Systems*

Designing AWS Data Lake for UK retailer.

**Technologies:** AWS S3, Spark, Kafka, Hive

*Feb-2017 - Nov 2017, Chief Software Engineer, EPAM Systems*

Designing and building a SQL driver for in-memory datagrid.

**Technologies:** Java, Apache Calcite, GigaSpaces XAP

*Jan-2016 - Feb-2017, Chief Software Engineer, EPAM Systems*

Designing and building a powerful Spark distribution <http://insightedge.io> from ground up. Allows to converge transactional and analytical processing within a single platform.

**Technologies:** Scala, Spark, GigaSpaces XAP

*Apr-2015 - Jan-2016, Lead Software Engineer, EPAM Systems*

Design and initial development of the next generation platform to optimize flight delays for American Airlines. Built on top of GigaSpaces XAP in-memory datagrid.

**Technologies:** Java, GigaSpaces XAP

*Apr-2015 - Aug-2015, Lead Software Engineer, EPAM Systems*

Design and initial development of Global Application Session Platform for QVC. This includes sharing HTTP and global session data as well as realtime and offline analytics. As part of this platform designed and developed Spring Session on top of GigaSpaces XAP.

**Technologies:** Java, GigaSpaces XAP, Spring, Spring Session, Netflix Hystrix

*Jan-2015 - Apr-2015, Lead Software Engineer, EPAM Systems*

Spring Data implementation for GigaSpaces XAP with a number of extensions. See <http://docs.gigaspace.com/sbp/spring-data.html>

**Technologies:** Java, Spring Data, GigaSpaces XAP

*Dec-2014 - Feb-2015, Lead Software Engineer, EPAM Systems*

Spark Streaming with GigaSpaces XAP integration library and reference example. See <http://docs.gigaspaces.com/sbp/spark-integration.html>

**Technologies:** Scala, Spark Streaming, GigaSpaces XAP

*Jun-2014 - Dec-2014, Lead Software Engineer, EPAM Systems*

Redesigning QVC's product catalog, availability and pricing services to address scalability and latency concerns. Initial development and consulting. The key of a new design was a move of product data (up to 1TB) from DB2 to in-memory data grid (GigaSpaces XAP). Pricing and availability computations leveraged data locality. As a result, the average response time reduced in 4 times with ability to scale almost linearly.

**Technologies:** Java, GigaSpaces XAP, Spring

*Feb-2014 - May-2014, Lead Software Engineer, EPAM Systems*

Apache Storm with GigaSpaces XAP integration as a reusable library. Reference architecture and best practices. See <http://docs.gigaspaces.com/sbp/storm-integration.html>

**Technologies:** Java, Apache Storm, GigaSpaces XAP, GigaSpaces Cloudify

*Jan-2014 - Feb-2014, Lead Software Engineer, EPAM Systems*

GigaSpaces XAP with Apache Kafka integration library. See <http://docs.gigaspaces.com/sbp/kafka-integration.html>

**Technologies:** Java, GigaSpaces XAP, Apache Kafka

*Oct-2013 - Jan-2014, Lead Software Engineer, EPAM Systems*

Gigaspace XAP proof of concept. OAuth2-like authentication service with HMAC authentication

**Technologies:** MongoDB, MySQL, Memcached, Java, GigaSpaces XAP, Spring, Jmeter, Epam Cloud

*Jun-2013 - Oct-2013, Lead Software Engineer, EPAM Systems*

Overmind. Scalable, high available and fault tolerant authentication service for Pivotlink products.

**Technologies:** MongoDB, Java, Scala, Memcached, Cloudera Hadoop stack, Spring, Play, jQuery, CoffeeScript

*Mar-2012 - Oct-2013, Lead Software Engineer, EPAM Systems*

Big Data analytics application designed for retail and consumer products companies looking to improve their operational decisions.

**Technologies:** PostgreSQL, Java, Memcached, Spring, Hibernate, Quartz, SOAP, Flex, C, Perl, keepalived, LVS

*Feb-2011 - Mar-2012, Lead Software Engineer, EPAM Systems*

IssueNet is a unique community of member banks, linked via a technology platform that gives syndicate to syndicate connectivity across the market.

**Technologies:** Oracle, EJB 2.1, Struts, JAXB, JBoss AS, jQuery

*Apr-2008 - Feb-2011, Senior Software Engineer, EPAM Systems*

IssueBook is a comprehensive web-based bookbuilding platform for managing the entire fixed-income new-issue process

**Technologies:** Oracle, Sybase, EJB 2.1, Spring, Hibernate, Struts, DWR, JBoss AS, Weblogic AS

*Jul-2009 - Aug-2010, OpenSource community*

Slag-gun is a 2D multiplayer flash game

**Technologies:** Java, Flash, ActionScript, Haskell

*Oct-2007 - Apr-2008, Software Engineer, EPAM Systems*

Vienna is a web application for business modeling which uses DB of the legacy product.

**Technologies:** Oracle, Spring, AOP, JDBC, POI

*May-2007 - Sep-2007, Software Engineer, Softline*

Web-based application for agricultural management and resource planning.

**Technologies:** Spring, Hibernate, JSF, Jboss AS

*Jun-2007 - Jul-2007, Software Engineer, Self employment*

Reporter Engine. Web-based application that is intended for articles publication and reporters' management. Allows to publish articles from mobile devices and web interface.

**Technologies:** MySQL, Spring, Hibernate, Tomcat

*May-2007 - Jul-2007, Software Engineer, Softline*

Business Reporting Mondrian Enhancements. Web-based application that is intended for generation business reports from data that is provided by OLAP services.

**Technologies:** PostgreSQL, OLAP, Mondrian, JPivot, Struts, Jboss AS

*Mar-2007 - May-2007, Software Engineer, SoftLine*

Exposure Model Baseline. Console tool that is intended for making prediction relying on some data of unknown nature.

**Technologies:** Neuron networks, XML, XSD, JAXB

*Sep-2006 - May-2007, Software Engineer, self employment*

OpenFaces Forum. Web-based application for holding discussions and posting user-generated content.

**Technologies:** MySQL, Spring, Struts, MySQL, Tomcat

## EDUCATION

Master, Computer Science, 2008  
Kharkiv National University of Radioelectronics

## CERTIFICATES

[Principles of Reactive Programming](#), Coursera  
[Functional Programming Principles in Scala](#), Coursera  
[Algorithms: Design and Analysis, Part 1](#), Coursera  
[Algorithms: Design and Analysis, Part 2](#), Coursera

## TALKS

[How we built InsightEdge](#), JavaDay 2016 (Devoxx)  
[Monads and Monoids: from daily Java to Big Data analytics in Scala](#), JavaDay 2015 (Devoxx)  
[Resilient distributed systems with Netflix Hystrix](#), Epam meetup  
[Understanding Parser Combinators in Scala](#), Epam meetup  
[Real-Time Big Data with Storm, Kafka and GigaSpaces](#), Epam meetup