

Kevin Liew

San Francisco | kliew@apache.org

Experience

Tradeshift, *Software Engineer*

Jun 2019 – present

- Converted the Supply Chain Finance onboarding program from a Spring-integrated pipeline into a RabbitMQ pipeline, to be used as a basis for further integration of messaging queues

Medallia, *Software Engineer II*

Jun 2018 – May 2019

- Supporting the deployment of a Text Analytics service to enable clients to drive actionable insights from topic and sentiment tagging of customer correspondence
- Logging telemetry with InfluxDB storage, Telegraf alerting, Grafana dashboards

Adenda Media, (*contractor*) *Backend-Server Engineer*

Jul 2016 – Jun 2017

- Achieved a significant boost in max throughput at Adenda's principal advertisement service by implementing a lambda architecture on Redis, Kafka, protobuf, Akka Streams, and MongoDB

Magnitude - Simba Technologies, *Software Engineer*

May 2015 – Mar 2017

- Worked in an agile team environment to develop and maintain Simba's Hive, Impala, Spark, Phoenix, and MongoDB ODBC drivers for Windows, Linux, and OSX
- Reverse engineered the Avatica protobuf-protocol to deliver the first Phoenix ODBC driver
- Implemented ODBC interfaces, Kerberos authentication, performance optimizations, and additional enterprise support-engineering to fortify Simba's offerings in a competitive market
- Deployed Hortonworks HDP Hive and Phoenix clusters, containerized MongoDB clusters for continuous integration testing to ensure client compatibility across server versions

Polycom, *Automation Engineer (co-op)*

Jan 2012 – Aug 2012

- Enhancing Polycom's test framework for VoIP, touchscreen, and video conferencing phones
-

Education

Simon Fraser University, B.ASc in Computer Engineering

2011 – 2016

Open Source Contributions

Apache Calcite, Data Management Framework

2017 – 2017

- Extending Calcite's SQL parser to support modifiable views and NoSQL dynamic-columns

Apache Hive, Data Warehouse on Hadoop Filesystem

2016 – 2017

- Designed a plugin framework and sample plugin to compress map-reduce results

Apache Phoenix, Distributed RDBMS on HBase

2015 – 2017

- Contributed to the integration of Calcite's SQL parser and cost-based query optimizer

GitHub Projects (<https://github.com/kliewkliew>)

Home Rental Cost, Web Scraper and Machine Learning Project

2018 – 2018

- Applied gradient boosting of regressions to train a machine learning model to estimate costs

Salad, Cornucopia, Redis Client and Cluster Management Framework

2016 – 2017

- Augmented the 'lettuce' Java 8 client for Scala applications
 - Automated management and deployment of Redis Cluster for horizontal scaling
-

Technical Skills

- Java 8, Scala, C++, Groovy
- Kafka, Redis, RabbitMQ
- Service-oriented architecture
- Spring, Akka, Play
- MongoDB, Calcite, Hive
- Thrift, Protocol Buffers